

CS207 Final Project: JBParallel

Jonah Kallenbach and Brendan Bozorgmir
jkallenbach@college.harvard.edu and bbozorgmir@college.harvard.edu

December 15, 2014

ALL CODE IS HERE: <https://github.com/jonahkall/JBParallel>

1 Intro, Setup and Installation

Our goal for this project was to implement a relatively simply but easily usable and fast library for parallelizing code. Until the very end of the project, we did not look at any parallel libraries, including OpenMP: at Cris's recommendation, we wanted to learn by implementing everything ourselves. The end result, thousands of lines of code and dozens of hours testing later is a success: we have sped up tons of code and had a lot of fun thinking and implementing algorithms in parallel, and other people, we hear, are using our library. In fact, by sacrificing a bit of generality, we were able to make our code faster than OpenMP's own functions (e.g. `omp parallel for`)!

We will first provide a quick guide to using our software on one's own computer. We tried to make the interface as simple and useful as possible.

Determine if your problem can be parallelized.

- Do you have a single for-loop which is extremely time-consuming?
- Are there no data dependencies, and does order not matter?
 - For example, a for loop to compute the fibonacci sequence would not work, because later iterations' values are dependent on earlier values.
- Are you sure this for loop is the slow part? For example, we noticed that on a lot of graph visualization problems, even though we were speeding up symplectic Euler step 3x, it didn't matter because SDLViewer rendering was actually the slow part.
- Is there enough work being done at each iteration, and is the loop itself big enough to justify parallelism? Keep in mind that parallelism requires significant overhead.

Endow your VM with multiple cores

- Power off your VM
- Go to settings for your CS207 VM
- Switch to 64 bit Ubuntu (optional)
- Under System → Motherboard, enable I/O APIC
- Enable hardware virtualization under System→ Virtualization
- Go to System→ Processor, and add 4 cores!
- Check by typing `lscpu` at your terminal that you indeed have 4 usable cores.
- Include our Library
- Download "jb_parallel.hpp" to your username-cs207 directory
- And `pctest.cpp` if you want to check that everything is working
- Include our library in your code:

```
#include "jb_parallel.hpp"
using namespace jb_parallel;
```

- Switch from clang to g++, and add `-fopenmp`
- In your makefile, switch from clang to g++

```
CXX := $(shell which g++) -std=c++11
```

- Also link in openmp:

```
CXXFLAGS += -O3 -funroll-loops -fopenmp -W -Wall -Wextra #-Wfatal-errors
```

- Write a unary function as a functor to do whatever was being done in your loop body
- If you don't care about the function's return value, use `for_each`, otherwise use `parallel_transform`
- For example, the unary function below does the position modification from symplectic Euler step from HW2.

```

template <typename F>
struct position_mod {
    double dt;
    void operator () (Node n) {
        n.position() += n.value().velocity * dt;
    }
    position_mod(double dt1) : dt(dt1) {};
};

```

- Turn your loop into a parallel `for_each` loop
- Maybe check first that you can do it as `std::for_each`
- Then use our parallel `for_each` (`jb_parallel::for_each`, or just `for_each` if you're already in our namespace) as:

```

position_mod<F> pm(dt);
for_each(g.node_begin(), g.node_end(), pm);

```

Enjoy your lightning fast code!

2 Code Provided

Our library can be found in `jb_parallel.hpp`, in which we've implemented several fundamental algorithms in parallel.

Here is a complete list of all algorithms implemented in that file

- `parallel_sort`, a function that sorts any two iterator values in a range by using `std::sort` and combining the results (currently only works on 4-core machines)
- `parallel_min`, a function that finds the min of any two iterator values in a range
- `parallel_transform` and `parallel_for_each`, which both apply a function to a range, with `for_each` modifying the range and `parallel_transform` returning a new one
- `parallel_reduce`, which applies a function to a range and then sums it (similar to `std::accumulate`, but with the added parameter of a function)

We've already sped up a lot of previous homework code and the extensions provided by our peers. The code that we have parallelized includes:

- `shallow_water.cpp` from HW4
- `mass_spring.cpp` from HW2

- `shallow_water_extension.cpp`, which is an extension that modifies HW4 to include a boat, written by Wenshuai Ye and Yuhao Zhu
- `kmeans.cpp`, which uses K-means clustering in parallel to help to color different clusters of any graph. Run this as follows:

```
make kmeans
./kmeans large_clustering_problem2.nodes 1
```

1 turns parallelization on, 0 turns it off. We noticed that we could actually make this a bit faster using raw OpenMP directives, but we did not want to sacrifice generality, and this provides a good example of us using our own library.

- `mesh_mass_spring.cpp` From George Lok's project
- `project2.cpp` from Brian Zhang's project

3 How fast is our code?

Parallelizing code can make it much faster (2-3 times for 4 core machines like the ones that we were using), but this comes at the cost of considerable overhead. Thus, it is important to consider the scale of your problem when deciding whether or not to use our algorithms. Generally speaking, these `jb_parallel` algorithms will lose to the comparable algorithm provided by standard until there are 10^6 entries to be performed on it. On a vector of 10^8 doubles, our sort algorithm was 2.11x as fast as `std`, `for_each` was 2.05x as fast, our min algorithm was 2.27x as fast, and our reduce algorithm was 2.5x as fast. Please consult the following figures for a more detailed viewing of our algorithms compared to `std` at different range values:

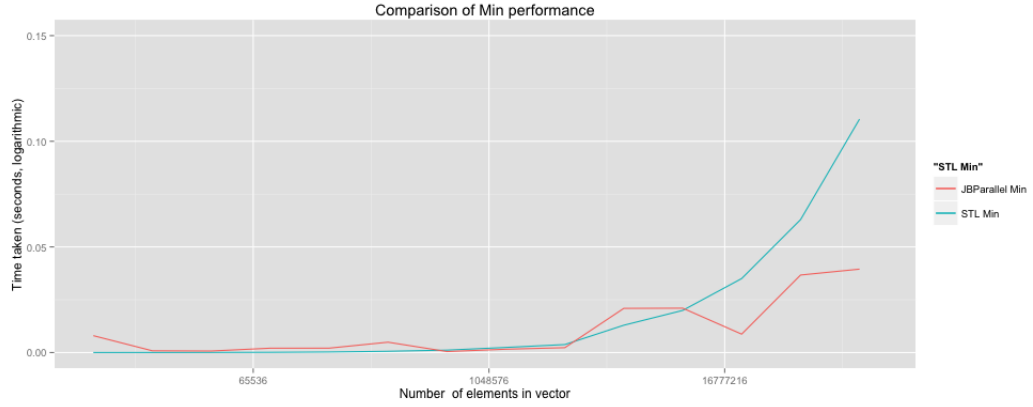


Figure 1: **Performance on min_element:**

We see that as the number of elements increases, our code begins. Fluctuations likely due to randomness (code too slow to run for a lot of trials)

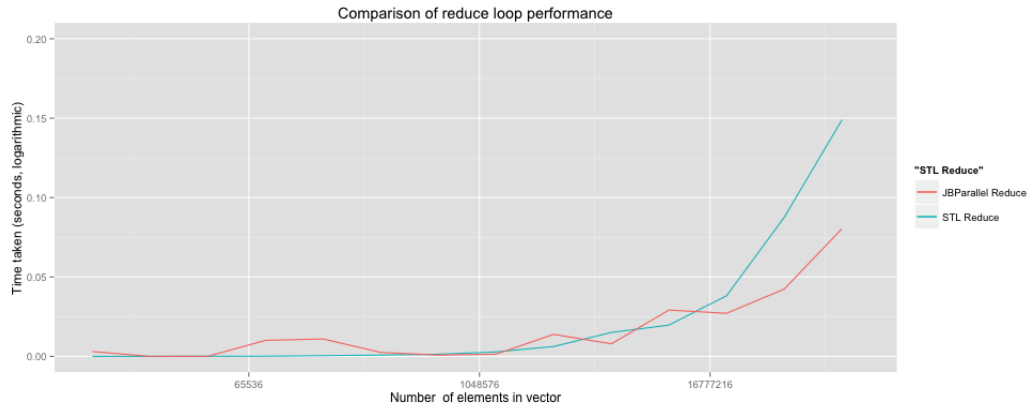


Figure 2: **Performance on reduce:**

Our code begins to outperform a serial reduction loop at about the 4×10^6 element mark, up to an ultimate speedup of about 2.0x. Same causes for randomness as above.

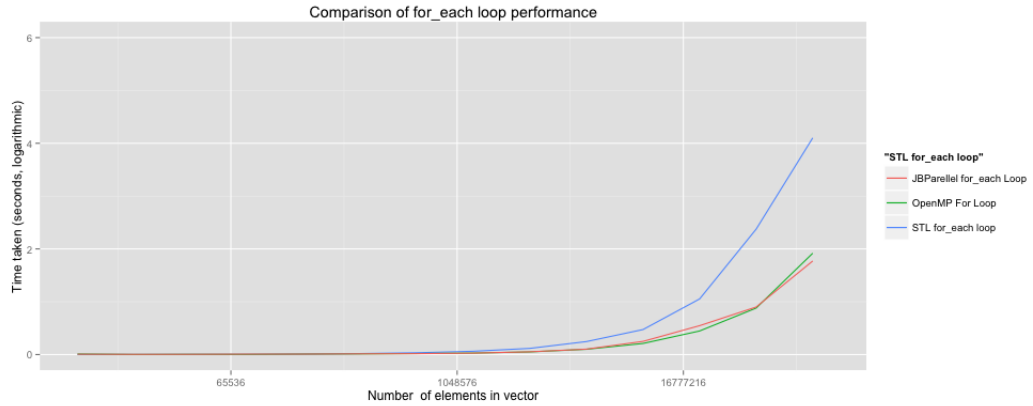


Figure 3: **Performance on for_each:**

Code begins to outperform at around the million element mark. It also consistently outperforms OpenMP on large problem sizes by about 1.2x. Same causes for randomness as above.

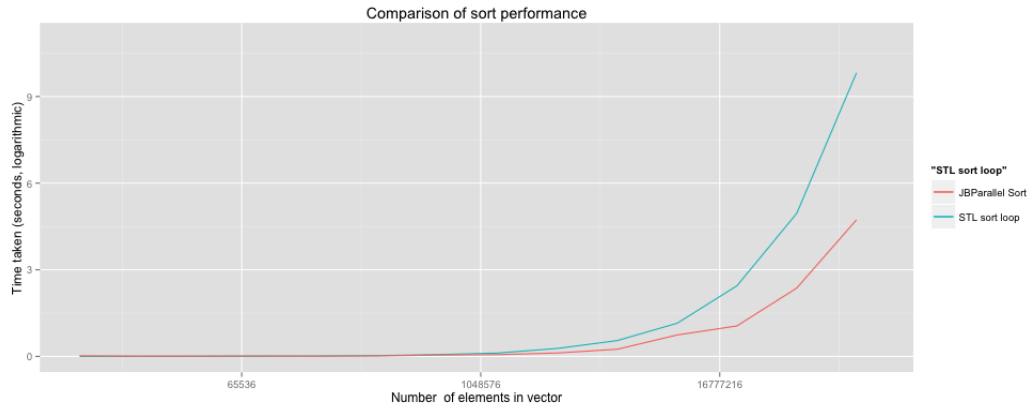


Figure 4: **Performance on sort:**

Again we win at large sizes, same causes for randomness as above.

It's important to note that in each of these figures, very little work is being done on each iteration of the parallel/non-parallel loop. The more work done on each iteration, however, the more the parallel algorithm wins, so in general practice, one would expect more speedup than is shown here. Other interesting results we came across while working on things in a test environment: nesting a parallel loop within non-parallel loops is very, very costly (unfortunately this is inevitable on the *step functions because we can't really compute things in the future), and using shared variables via OpenMP's **shared** keyword.

Here is a typical run of our testing file, `pctest.cpp`, which is in our project directory.

```
Min element found by serial: 0
Serial Min: 0.0878805s
Min element found by parallel: 0
jb_parallel parallel_min: 0.0462203s
```

```
Serial Parallel Loop: 2.64941s
OMP Parallel Loop: 1.05172s
JB Parallel_Transform: 0.480762s
```

```
Serial Sort: 2.30043s
Parallel Sort: 0.930822s
```

```
Serial reduction: 0.0699698s
0
Parallel reduction: 0.0174481s
0
```

To run this inside our repo and get similar results, simply run `make ptest`, followed by `./ptest`.

As you can see, this run was slightly atypical in terms of speed, but in line with the numbers stated above. We found that speed data was difficult to collect consistently, so we wrote `speedtest.py` to allow us to run our code easily hundreds or thousands of times (typically we would see acceleration after multiple runs, probably due to some sort of caching issue). Of course the issue there was that in order to see speedup, we needed to be working with large amounts of data, but then the task would take long enough that running the code say, a thousand times would become impractical.

Another tool which assisted us enormously in our quest to make code fast was `valgrind`'s `callgrind` profiler tool, and `kcachegrind`, a tool which provides a GUI to visualize `callgrind` output. For example, in speeding up George's code (they used our extension and sped their code up, and we used our extension and sped their code up in parallel :)), we first ran:

```
valgrind --tool=callgrind ./mesh_mass_spring data/sphere2.nodes data/sphere2.tris
```

The program `kcachegrind` allows one to see the output of this profiler using a nice GUI which allows for visualization of the call graph, call map, and raw call counts for all functions, which is extremely helpful in profiling. We found in George's code that the `mesh_shape_volume` function was particularly slow, and was also a great use case of `parallel_reducer`, so we used this to achieve about 2x speed up (this will only work with very large inputs).

4 Collaboration Evaluations

The extensions provided by other students on Piazza was invaluable to our project. We were given large amounts of very well written C++ code that included `for_loops` that could be easily parallelized using our `jb_parallel` library. Note that for all of these, we only looked at a specific subset of the code in the hopes of optimizing it. Also please take any criticisms we have with a grain of salt, as we often grabbed code from these projects while they were still in an intermediate state, and this probably accounts for any problems we had.

4.1 George Lok and Serguei Balanovich

Very well written, understandable code. We really liked the way their code was laid out. This allowed us to quickly, and modularly, modify their code. And we understand they used our package as well, independently of us!

4.2 Brian Zhang and Tarik Moon

At the time we received it, the code wasn't super well documented, lots of stuff un-commented and lots of stuff commented out. Again, we think we just received the code at an unfavorable time, when that group was still figuring stuff out (because we could pull after any commit). Nonetheless, we were extremely impressed with the results of the project.

4.3 Wenshuai Ye and Yuhao Zhu

Wenshuai and Yuhao wrote a great extension to `shallow_water` called `shallow_water_extension.cpp`. We also used their `Mesh.cpp` file, which we renamed `Mesh2.cpp` to not conflict with our existing code.

In `Mesh2.cpp`, we rewrote their `triangle_iterator` to make it a Random Access Iterator.

In `shallow_water_extension.cpp`, we refactored the code inside of the `hyperbolic_step` to create a functor called `FluxUpdater` with the exact same logic as before. This allowed us to use their code in our `jb_parallel::for_each`, which parallelized their code.